

# The Application Layer

Michael Brodskiy

Professor: E. Bernal Mor

September 25, 2023

- Network Applications
  - Social networking
  - Web
  - Text messaging
  - E-mail
  - Multi-user network games
  - Streaming stored video (YouTube, Hulu, Netflix)
  - P2P File Sharing
  - And many more
- Creating Network Applications
  - Write programs that:
    - \* Run on (different) end systems
    - \* Communicate over network
    - \* For example, web server software communicates with browser software
  - No need to write software for network-core devices (intermediate nodes)
    - \* Network-core devices do not run user applications
    - \* Applications on end systems allow for rapid application development and propagation
- Application Architecture
  - Network architecture — a set of layers and protocols
    - \* It is fixed, and provides the network application developer with specific set of services

- Application Architecture — define how the application is structured over various end systems
  - \* Designed by the application developer
  - \* Predominant architectural paradigms
    - Client-server
    - Peer-to-peer (P2P)
- Client-server Architecture
  - Server
    - \* Always-on host
    - \* Permanent IP-address (like ID)
    - \* Often in data centers, for scaling
  - Clients
    - \* Contact, communicate with server
    - \* May be intermittently connected
    - \* May have dynamic IP addresses
    - \* Do not communicate directly with each other
  - Examples: HTTP, IMAP, SFTP
- Peer-Peer (P2P) Architecture
  - No always-on server
  - Arbitrary end systems directly communicate
  - Peers request service from other peers, provide service in return to other peers
    - \* Self scalability — new peers bring new service capacity, as well as new service demands
  - Peers are intermittently connected and change IP addresses
    - \* Complex management
  - Example: P2P File Sharing
- Process Communication
  - Process — program running within a host
    - \* Within same host, two processes communicate using inter-process communication, defined by OS (Operating System)
    - \* Processes in different hosts communicate by exchanging messages
  - Client process — process that initiates communication
  - Server process — process that wants to be contacted

- Note: applications with P2P architectures have client processes & server processes
- Sockets
  - Process send/receives messages to/from its socket
  - Socket analogous to door
    - \* Sending process shoves message out the door
    - \* Sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
- Addressing Processes
  - To receive messages, a process must have an identifier
  - Host device has a unique IP address
  - Identifier includes both IP address and port numbers associated with process on host
    - \* HTTP server: 80
    - \* Mail server: 25
  - To send HTTP message to gaia.cs.umass.edu web server:
    - \* IP address: 128.119.245.12
    - \* Port number: 80
- An Application Layer Protocol Defines:
  - Types of messages exchanged
    - \* Example: request, response
  - Message syntax
    - \* What fields in messages & how fields are delineated
  - Message semantics
    - \* Meaning of information in fields
  - Rules for when and how processes send & respond to messages
- Application Layer Protocols can be
  - Open protocols
    - \* Defined in RFCs, everyone has access to protocol definition
    - \* Allows for interoperability
    - \* Example: HTTP, SMTP
  - Proprietary protocols

- \* Example: Skype
- Transport Layer Services for Applications
  - Transport layer is on the other side of the “door”
  - There are multiple Transport-layer protocols that provide different services
  - The application developer must choose a Transport-layer protocol, depending on the services needed by the application
    - \* Examples: priority mail, express mail, certified mail
  - A Transport-layer protocol can provide a different array of services
- Transport Services
  - Data integrity/reliable transport
    - \* Some apps (*e.g.* file transfer, web transactions) require 100% reliable data transfer
    - \* Other apps (*e.g.* audio) can tolerate some loss
  - Timing
    - \* Some apps (*e.g.* Internet telephony, interactive games) require low delay to be “effective”
  - Throughput
    - \* Some apps (*e.g.*, multimedia) require minimum amount of throughput to be “effective”
    - \* Other apps (“elastic apps”) make use of whatever throughput they get
  - Security
    - \* Encryption, data integrity, ...
- Internet Transport Protocol Services
  - TCP Service
    - \* Reliable transport between sending and receiving processes
    - \* Flow control — sender will not overwhelm receiver
    - \* Congestion control — throttle sender when network overloaded
    - \* Does not provide timing, minimum throughput guarantee, security
    - \* Connection-oriented service: setup required between client and service processes
  - UDP Service:
    - \* Unreliable data transfer between sending and receiving process
    - \* Does not provide reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup

- \* Connectionless service: no setup required
- Vanilla TCP & UDP sockets
  - No encryption
  - Clear text passwords sent into socket traverse Internet in clear text
- Transport Layer Security (TLS)
  - Provides encrypted TCP connections
  - Data integrity
  - End-point authentication
  - TSL implemented in Application Layer
    - \* Applications use TLS libraries, that use TCP in turn
  - TLS socket API
    - \* Clear text sent into socket traverse Internet encrypted
  - Datagram Transport Layer Service (DTLS) protocol
    - \* Adaptation of TLS to run over connectionless protocols such as UDP
- Designing Network Applications
  - It is a complex process
  - Requires knowledge of programming, software engineering, and networking
  - From a networking point of view, there are two major decisions:
    1. Type of application (aka Application Architecture)
      - \* Client-server vs. peer-to-peer
    2. Services requested to the Transport Layer
      - \* *E.g.* reliable vs. unreliable data transfer
- Web and HTTP
  - Web page consists of several objects, each of which can be store on different Web servers
  - Object can be HTML, JPEG, Java applet, audio file, etc.
  - Web page consists of a base HTML-file, which includes several referenced objects
  - Each object is addressable by a URL (Uniform Resource Locator), *e.g.*,

$$\underbrace{\text{www.someschool.edu}}_{\text{host name}} / \underbrace{\text{someDept/pic.gif}}_{\text{path name}}$$

- HTTP Overview
  - HTTP — Hypertext Transfer Protocol
  - Web’s application layer protocol
  - Client/server model
    - \* Client: Browser that requests, receives (using HTTP protocol) and “displays” Web objects
    - \* Server: Web server sends (using HTTP protocol) objects in response to requests
  - Versions
    - \* HTTP/1.0 (RFC1945)
      - Original HTTP version (early 1990s)
    - \* HTTP/1.1 (RFC7230,...)
      - Used by most of the HTTP transactions
    - \* HTTP/2 (RFC7540,...)
      - Standardized in 2015 and increasingly used by browsers and servers
    - \* HTTP/3 (RFC9114)
      - IETF published it as a proposed standard in June 2022
  - HTTP used TCP (except HTTP/3):
    - \* Client initiates TCP connect (creates socket) to server, port 80
    - \* Server accepts TCP connection from client
    - \* HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
    - \* TCP connection closed
  - HTTP is “stateless”
    - \* Server maintains no information about past client requests
  - Protocols that maintain “state” are complex
    - \* Past history (state) must be maintained
    - \* If server/client crashes, their view of “state” may be inconsistent, must be reconciled
  - Non-persistent HTTP
    1. TCP connection opened
    2. At most one object sent over TCP connection
    3. TCP connection closed
    - \* Downloading multiple objects requires multiple connections
  - Persistent HTTP

1. TCP connection opened to a server
  2. Multiple objects can be sent over single TCP connection between client, and that server
  3. TCP connection closed
- RTT (Round-Trip Time) definition: time for a small packet to travel from client to server and back
    - \* Includes propagation, queueing, and processing delays
  - Non-persistent HTTP response time (per object)
    - \* One RTT to initiate TCP connection
    - \* One RTT for HTTP request/response
    - \* File/object transmission time
    - \* Non-persistent HTTP time  $\approx 2\text{RTT} + \text{transmission time}$  (per object)
  - Non-persistent HTTP issues:
    - \* Requires 2 RTTs per object
    - \* OS overhead for each TCP connection
    - \* Browsers often open parallel TCP connections to fetch referenced objects in parallel
  - Persistent HTTP/1.1:
    - \* Server leaves connection open after sending response
    - \* Subsequent HTTP messages between same client/server sent over open connection
    - \* Client sends requests as soon as it encounters a referenced object
    - \* As little as one RTT for all the referenced objects (cutting response time in half)
  - HTTP Request Message
    - \* Two types of HTTP messages: request, response
    - \* HTTP request message is in ASCII (human-readable format)
- HTTP Methods
    - GET method
      - \* To request an object
      - \* If user data (like form input), the data is included in URL field of HTTP GET request message (following a '?'):  
`www.somesite.com/animalsearch?table&chair`
    - POST method
      - \* Web page often includes form input

- \* User input sent from client to server in entity body of HTTP POST request message
- HEAD method
  - \* Requests headers (only) that would be returned if specified URL were requested with an HTTP GET method
- PUT method
  - \* Uploads new file (object) to server
  - \* Completely replaces file that exists at specified URL with content in entity body
- DELETE method
  - \* Deletes file specified in the URL field
- HTTP Response Status Codes
  - 200 — OK
    - \* Request succeeded, requested object later in this message
  - 301 — Moved Permanently
    - \* Requested object moved, new location specified later in this msg (Location:)
  - 400 — Bad Request
    - \* Request msg not understood by server
  - 404 — Not Found
    - \* Requested document not found on this server
  - 505 HTTP — Version Not Supported
- HTTP/2
  - Key goal: decreased delay in multiple-object HTTP requests
  - HTTP/1.1: introduced multiple, pipelined GETs over single TCP connection (persistent HTTP)
    - \* FCFS (First-Come-First-Served) scheduling: server responds in-order to GET requests
    - \* Head-Of-Line (HOL) blocking: with FCFS, small object may have to wait for transmission behind large object(s)
    - \* Loss recovery (retransmitting lost TCP segments) stalls object transmission
  - HTTP/2: increased flexibility at server in sending objects to client
    - \* Methods, status codes, most header fields unchanged from HTTP 1.1
    - \* Transmission order of requested objects based on client-specified object priority (not necessarily FCFS)



- \* Push unrequested objects to client
  - \* Divide objects into frames, schedule frames to mitigate HOL blocking
- HTTP/2 to HTTP/3
  - HTTP/2 over single TCP connection means:
    - \* Recovery from packet loss still stalls all object transmissions
      - As in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling and increase overall application data throughput
    - \* No security over vanilla TCP connection
  - HTTP/3 operates over QUIC, a transport protocol built on UDP
    - \* Adds security, per object error, and congestion-control (more pipelining) over UDP
- Maintaining User/Server State: Cookies
  - Recall: HTTP GET/response interaction is stateless
  - Web sites and client browser use cookies to maintain some state between transactions
  - Four components:
    1. Cookie header line of HTTP response message
    2. Cookie header line in next HTTP request message
    3. Cookie file kept on user's host, managed by user's browser
    4. Back-end database at Web site
- HTTP Cookies: Comments
  - What cookies can be used for:
    - \* Authorization
    - \* Shopping carts
    - \* Recommendations
    - \* User session state (Web e-mail)
  - Challenge: How to keep state
    - \* Protocol endpoints: maintain state at sender/receiver over multiple transactions
    - \* Cookies: HTTP messages carry state
  - Cookies and Privacy:
    - \* Cookies permit sites to learn a lot about you on their site

- \* Third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites
- Web Caches (Proxy Server)
  - Goal: satisfy client request without involving origin server
  - User configures browser to point to a web cache
  - Browser sends all HTTP requests to cache
    - \* If object in cache, cache returns object to client
    - \* Else cache requests object from origin server, caches received object, then returns object to client
  - Web cache acts as both client and server
    - \* Server for original requesting client
    - \* Client to origin server
  - Typically, cache is installed by ISP (university, company, residential ISP)
  - Why Web Caching?
    - \* Reduce response time for client request
      - Cache is closer to client
    - \* Reduce traffic on an institution's access link
    - \* Internet is dense with caches
      - Enables “poor” content providers to more effectively deliver content
- Conditional GET
  - Web cache: specify date of cached copy in HTTP request:
 

**If-modified-since: <date>**
  - Server: response contains no object if cached copy is up-to-date
 

**HTTP/1.1 304 Not Modified**
  - Goal: don't send object if cache has up-to-date cached version
    - \* No object transmission delay
    - \* Lower link utilization
- E-Mail
  - Three major components
    1. User agents (UA)
      - \* A.k.a. “mail reader”

- \* Composing, editing, reading mail messages
  - \* E.g., Outlook, iPhone mail client
  - \* Outgoing, incoming messages stored on server
- 2. Mail servers
  - \* Mailbox contains incoming messages for user
  - \* Message queue of outgoing (to be sent) mail messages
- 3. Simple Mail Transfer Protocol (SMTP)
  - \* Between mail servers to send e-mail messages
    - Client: Sending mail server
    - Server: Receiving mail server
  - \* SMTP can also be used by user agents to send e-mails to mail server
- E-Mail: SMTP (RFC 5321)
  - Uses TCP to reliably transfer e-mail message from client to server (mail server port 25)
  - Direct transfer: sending mail server (acting like client) to receiving mail server
  - After TCP connection is established, three phases of transfer
    - \* Handshaking (greeting)
    - \* Transfer of messages
    - \* Closure
  - Command/response interaction (like HTTP)
    - \* Commands: ASCII text
    - \* Response: status code and phrase
  - Messages must be in 7-bit ASCII
  - Comparison with HTTP:
    - \* HTTP: pull; SMTP: push
    - \* Both have ASCII command/response interaction and status codes
    - \* HTTP: each object encapsulated in its own response message; SMTP: multiple objects sent in multipart message
    - \* SMTP: multiple objects sent in multipart message
    - \* SMTP uses persistent connections
    - \* SMTP requires message (header & body) to be in 7-bit ASCII
- E-Mail Access Protocols
  - SMTP: delivery/storage of e-mail messages to receiver's server
  - E-Mail access protocol: Bob's user agent starts the communication for retrieval from server

- \* IMAP: Internet Mail Access Protocol (RFC 3501): messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
  - HTTP: Gmail, Hotmail, etc. provides web-based interface to send or to retrieve e-mail messages
- Domain Name System (DNS)
  - People have many identifiers:
    - \* SSN, name, passport number, etc.
  - Internet hosts, routers
    - \* IP addresss — used for addressing packets
    - \* Name, like CS.UMASS.EDU is used by humans
  - Distributed database implemented in hierarchy of many name servers
  - Application-layer protocol: hosts and name servers communicate to resolve names (address/name translation)
    - \* Note: core Internet function, implemented as application-layer protocol
    - \* Complexity at network's “edge”