

The Application Layer

Michael Brodskiy

Professor: E. Bernal Mor

September 28, 2023

- Network Applications
 - Social networking
 - Web
 - Text messaging
 - E-mail
 - Multi-user network games
 - Streaming stored video (YouTube, Hulu, Netflix)
 - P2P File Sharing
 - And many more
- Creating Network Applications
 - Write programs that:
 - * Run on (different) end systems
 - * Communicate over network
 - * For example, web server software communicates with browser software
 - No need to write software for network-core devices (intermediate nodes)
 - * Network-core devices do not run user applications
 - * Applications on end systems allow for rapid application development and propagation
- Application Architecture
 - Network architecture — a set of layers and protocols
 - * It is fixed, and provides the network application developer with specific set of services

- Application Architecture — define how the application is structured over various end systems
 - * Designed by the application developer
 - * Predominant architectural paradigms
 - Client-server
 - Peer-to-peer (P2P)
- Client-server Architecture
 - Server
 - * Always-on host
 - * Permanent IP-address (like ID)
 - * Often in data centers, for scaling
 - Clients
 - * Contact, communicate with server
 - * May be intermittently connected
 - * May have dynamic IP addresses
 - * Do not communicate directly with each other
 - Examples: HTTP, IMAP, SFTP
- Peer-Peer (P2P) Architecture
 - No always-on server
 - Arbitrary end systems directly communicate
 - Peers request service from other peers, provide service in return to other peers
 - * Self scalability — new peers bring new service capacity, as well as new service demands
 - Peers are intermittently connected and change IP addresses
 - * Complex management
 - Example: P2P File Sharing
- Process Communication
 - Process — program running within a host
 - * Within same host, two processes communicate using inter-process communication, defined by OS (Operating System)
 - * Processes in different hosts communicate by exchanging messages
 - Client process — process that initiates communication
 - Server process — process that wants to be contacted

- Note: applications with P2P architectures have client processes & server processes
- Sockets
 - Process send/receives messages to/from its socket
 - Socket analogous to door
 - * Sending process shoves message out the door
 - * Sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
- Addressing Processes
 - To receive messages, a process must have an identifier
 - Host device has a unique IP address
 - Identifier includes both IP address and port numbers associated with process on host
 - * HTTP server: 80
 - * Mail server: 25
 - To send HTTP message to gaia.cs.umass.edu web server:
 - * IP address: 128.119.245.12
 - * Port number: 80
- An Application Layer Protocol Defines:
 - Types of messages exchanged
 - * Example: request, response
 - Message syntax
 - * What fields in messages & how fields are delineated
 - Message semantics
 - * Meaning of information in fields
 - Rules for when and how processes send & respond to messages
- Application Layer Protocols can be
 - Open protocols
 - * Defined in RFCs, everyone has access to protocol definition
 - * Allows for interoperability
 - * Example: HTTP, SMTP
 - Proprietary protocols

- * Example: Skype
- Transport Layer Services for Applications
 - Transport layer is on the other side of the “door”
 - There are multiple Transport-layer protocols that provide different services
 - The application developer must choose a Transport-layer protocol, depending on the services needed by the application
 - * Examples: priority mail, express mail, certified mail
 - A Transport-layer protocol can provide a different array of services
- Transport Services
 - Data integrity/reliable transport
 - * Some apps (*e.g.* file transfer, web transactions) require 100% reliable data transfer
 - * Other apps (*e.g.* audio) can tolerate some loss
 - Timing
 - * Some apps (*e.g.* Internet telephony, interactive games) require low delay to be “effective”
 - Throughput
 - * Some apps (*e.g.*, multimedia) require minimum amount of throughput to be “effective”
 - * Other apps (“elastic apps”) make use of whatever throughput they get
 - Security
 - * Encryption, data integrity, ...
- Internet Transport Protocol Services
 - TCP Service
 - * Reliable transport between sending and receiving processes
 - * Flow control — sender will not overwhelm receiver
 - * Congestion control — throttle sender when network overloaded
 - * Does not provide timing, minimum throughput guarantee, security
 - * Connection-oriented service: setup required between client and service processes
 - UDP Service:
 - * Unreliable data transfer between sending and receiving process
 - * Does not provide reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup

- * Connectionless service: no setup required
- Vanilla TCP & UDP sockets
 - No encryption
 - Clear text passwords sent into socket traverse Internet in clear text
- Transport Layer Security (TLS)
 - Provides encrypted TCP connections
 - Data integrity
 - End-point authentication
 - TSL implemented in Application Layer
 - * Applications use TLS libraries, that use TCP in turn
 - TLS socket API
 - * Clear text sent into socket traverse Internet encrypted
 - Datagram Transport Layer Service (DTLS) protocol
 - * Adaptation of TLS to run over connectionless protocols such as UDP
- Designing Network Applications
 - It is a complex process
 - Requires knowledge of programming, software engineering, and networking
 - From a networking point of view, there are two major decisions:
 1. Type of application (aka Application Architecture)
 - * Client-server vs. peer-to-peer
 2. Services requested to the Transport Layer
 - * *E.g.* reliable vs. unreliable data transfer
- Web and HTTP
 - Web page consists of several objects, each of which can be store on different Web servers
 - Object can be HTML, JPEG, Java applet, audio file, etc.
 - Web page consists of a base HTML-file, which includes several referenced objects
 - Each object is addressable by a URL (Uniform Resource Locator), *e.g.*,

$$\underbrace{\text{www.someschool.edu}}_{\text{host name}}/\underbrace{\text{someDept/pic.gif}}_{\text{path name}}$$

- HTTP Overview
 - HTTP — Hypertext Transfer Protocol
 - Web’s application layer protocol
 - Client/server model
 - * Client: Browser that requests, receives (using HTTP protocol) and “displays” Web objects
 - * Server: Web server sends (using HTTP protocol) objects in response to requests
 - Versions
 - * HTTP/1.0 (RFC1945)
 - Original HTTP version (early 1990s)
 - * HTTP/1.1 (RFC7230,...)
 - Used by most of the HTTP transactions
 - * HTTP/2 (RFC7540,...)
 - Standardized in 2015 and increasingly used by browsers and servers
 - * HTTP/3 (RFC9114)
 - IETF published it as a proposed standard in June 2022
 - HTTP used TCP (except HTTP/3):
 - * Client initiates TCP connect (creates socket) to server, port 80
 - * Server accepts TCP connection from client
 - * HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
 - * TCP connection closed
 - HTTP is “stateless”
 - * Server maintains no information about past client requests
 - Protocols that maintain “state” are complex
 - * Past history (state) must be maintained
 - * If server/client crashes, their view of “state” may be inconsistent, must be reconciled
 - Non-persistent HTTP
 1. TCP connection opened
 2. At most one object sent over TCP connection
 3. TCP connection closed
 - * Downloading multiple objects requires multiple connections
 - Persistent HTTP

1. TCP connection opened to a server
 2. Multiple objects can be sent over single TCP connection between client, and that server
 3. TCP connection closed
- RTT (Round-Trip Time) definition: time for a small packet to travel from client to server and back
 - * Includes propagation, queueing, and processing delays
 - Non-persistent HTTP response time (per object)
 - * One RTT to initiate TCP connection
 - * One RTT for HTTP request/response
 - * File/object transmission time
 - * Non-persistent HTTP time $\approx 2\text{RTT} + \text{transmission time}$ (per object)
 - Non-persistent HTTP issues:
 - * Requires 2 RTTs per object
 - * OS overhead for each TCP connection
 - * Browsers often open parallel TCP connections to fetch referenced objects in parallel
 - Persistent HTTP/1.1:
 - * Server leaves connection open after sending response
 - * Subsequent HTTP messages between same client/server sent over open connection
 - * Client sends requests as soon as it encounters a referenced object
 - * As little as one RTT for all the referenced objects (cutting response time in half)
 - HTTP Request Message
 - * Two types of HTTP messages: request, response
 - * HTTP request message is in ASCII (human-readable format)
- HTTP Methods
 - GET method
 - * To request an object
 - * If user data (like form input), the data is included in URL field of HTTP GET request message (following a '?'):
`www.somesite.com/animalsearch?table&chair`
 - POST method
 - * Web page often includes form input

- * User input sent from client to server in entity body of HTTP POST request message
- HEAD method
 - * Requests headers (only) that would be returned if specified URL were requested with an HTTP GET method
- PUT method
 - * Uploads new file (object) to server
 - * Completely replaces file that exists at specified URL with content in entity body
- DELETE method
 - * Deletes file specified in the URL field
- HTTP Response Status Codes
 - 200 — OK
 - * Request succeeded, requested object later in this message
 - 301 — Moved Permanently
 - * Requested object moved, new location specified later in this msg (Location:)
 - 400 — Bad Request
 - * Request msg not understood by server
 - 404 — Not Found
 - * Requested document not found on this server
 - 505 HTTP — Version Not Supported
- HTTP/2
 - Key goal: decreased delay in multiple-object HTTP requests
 - HTTP/1.1: introduced multiple, pipelined GETs over single TCP connection (persistent HTTP)
 - * FCFS (First-Come-First-Served) scheduling: server responds in-order to GET requests
 - * Head-Of-Line (HOL) blocking: with FCFS, small object may have to wait for transmission behind large object(s)
 - * Loss recovery (retransmitting lost TCP segments) stalls object transmission
 - HTTP/2: increased flexibility at server in sending objects to client
 - * Methods, status codes, most header fields unchanged from HTTP 1.1
 - * Transmission order of requested objects based on client-specified object priority (not necessarily FCFS)

- * Push unrequested objects to client
 - * Divide objects into frames, schedule frames to mitigate HOL blocking
- HTTP/2 to HTTP/3
 - HTTP/2 over single TCP connection means:
 - * Recovery from packet loss still stalls all object transmissions
 - As in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling and increase overall application data throughput
 - * No security over vanilla TCP connection
 - HTTP/3 operates over QUIC, a transport protocol built on UDP
 - * Adds security, per object error, and congestion-control (more pipelining) over UDP
- Maintaining User/Server State: Cookies
 - Recall: HTTP GET/response interaction is stateless
 - Web sites and client browser use cookies to maintain some state between transactions
 - Four components:
 1. Cookie header line of HTTP response message
 2. Cookie header line in next HTTP request message
 3. Cookie file kept on user's host, managed by user's browser
 4. Back-end database at Web site
- HTTP Cookies: Comments
 - What cookies can be used for:
 - * Authorization
 - * Shopping carts
 - * Recommendations
 - * User session state (Web e-mail)
 - Challenge: How to keep state
 - * Protocol endpoints: maintain state at sender/receiver over multiple transactions
 - * Cookies: HTTP messages carry state
 - Cookies and Privacy:
 - * Cookies permit sites to learn a lot about you on their site

- * Third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites
- Web Caches (Proxy Server)
 - Goal: satisfy client request without involving origin server
 - User configures browser to point to a web cache
 - Browser sends all HTTP requests to cache
 - * If object in cache, cache returns object to client
 - * Else cache requests object from origin server, caches received object, then returns object to client
 - Web cache acts as both client and server
 - * Server for original requesting client
 - * Client to origin server
 - Typically, cache is installed by ISP (university, company, residential ISP)
 - Why Web Caching?
 - * Reduce response time for client request
 - Cache is closer to client
 - * Reduce traffic on an institution's access link
 - * Internet is dense with caches
 - Enables “poor” content providers to more effectively deliver content
- Conditional GET
 - Web cache: specify date of cached copy in HTTP request:

If-modified-since: <date>
 - Server: response contains no object if cached copy is up-to-date

HTTP/1.1 304 Not Modified
 - Goal: don't send object if cache has up-to-date cached version
 - * No object transmission delay
 - * Lower link utilization
- E-Mail
 - Three major components
 1. User agents (UA)
 - * A.k.a. “mail reader”

- * Composing, editing, reading mail messages
 - * E.g., Outlook, iPhone mail client
 - * Outgoing, incoming messages stored on server
- 2. Mail servers
 - * Mailbox contains incoming messages for user
 - * Message queue of outgoing (to be sent) mail messages
- 3. Simple Mail Transfer Protocol (SMTP)
 - * Between mail servers to send e-mail messages
 - Client: Sending mail server
 - Server: Receiving mail server
 - * SMTP can also be used by user agents to send e-mails to mail server
- E-Mail: SMTP (RFC 5321)
 - Uses TCP to reliably transfer e-mail message from client to server (mail server port 25)
 - Direct transfer: sending mail server (acting like client) to receiving mail server
 - After TCP connection is established, three phases of transfer
 - * Handshaking (greeting)
 - * Transfer of messages
 - * Closure
 - Command/response interaction (like HTTP)
 - * Commands: ASCII text
 - * Response: status code and phrase
 - Messages must be in 7-bit ASCII
 - Comparison with HTTP:
 - * HTTP: pull; SMTP: push
 - * Both have ASCII command/response interaction and status codes
 - * HTTP: each object encapsulated in its own response message; SMTP: multiple objects sent in multipart message
 - * SMTP: multiple objects sent in multipart message
 - * SMTP uses persistent connections
 - * SMTP requires message (header & body) to be in 7-bit ASCII
- E-Mail Access Protocols
 - SMTP: delivery/storage of e-mail messages to receiver's server
 - E-Mail access protocol: Bob's user agent starts the communication for retrieval from server

- * IMAP: Internet Mail Access Protocol (RFC 3501): messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
- HTTP: Gmail, Hotmail, etc. provides web-based interface to send or to retrieve e-mail messages
- Domain Name System (DNS)
 - People have many identifiers:
 - * SSN, name, passport number, etc.
 - Internet hosts, routers
 - * IP addresss — used for addressing packets
 - * Name, like CS.UMASS.EDU is used by humans
 - Distributed database implemented in hierarchy of many name servers
 - Application-layer protocol: hosts and name servers communicate to resolve names (address/name translation)
 - * Note: core Internet function, implemented as application-layer protocol
 - * Complexity at network's "edge"
- DNS Services
 - Hostname to IP address translation
 - Host aliasing
 - * Canonical, alias names
 - Mail server aliasing
 - Load distribution
 - * Replaced Web servers: many IP addresses correspond to one name
 - Why not centralize DNS?
 - * Single point of failure
 - * Traffic volume
 - * Distant centralized database
 - * Maintenance
 - * It doesn't scale!
 - Client wants IP address for Amazon: 1st approximation:
 - * Query sent to root server to find .com DNS server (root)
 - * Query sent to .com DNS server to get amazon.com DNS server (top level domain)
 - * Query sent to amazon.com DNS server to get IP address for www.amazon.com (authoritative)

- Local DNS Server

- Local DNS Server — helps hosts to resolve a query
- Does not strictly belong to hierarchy
- Each ISP or institution (residential ISP, company, university) has one
 - * Also called “default name server”
- When host makes DNS query, query is sent to its local DNS server
 - * Has local cache of recent name-to-address translation pairs (but may be out of date!)
 - * Acts as proxy, forwards query to hierarchy

- DNS: Root Name Servers

- Official, contact-of-last-resort by name servers that cannot resolve name
- Root name servers provide the IP addresss of TLD servers
- Incredibly important Internet function
 - * Internet couldn’t function without it!
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain
- 13 logical root name “servers” worldwide, each “server” replicated many times (about 200 servers in the US)

- TLD and Authoritative Servers

- Top-level domain (TLD) servers:
 - * Responsible for .com, .org, .net, .edu, etc.
 - * Network solutions maintains servers for .com, .net TLD
 - * Educause for .edu TLD
- Authoritative DNS servers:
 - * Organization’s own DNS server(s), providing authoritative hostname to IP mappings for organizations named hosts

- DNS Name Resolution: Iterated Query

- Example: host at ENGINEERING.NYU.EDU wants IP address for GAIA.CS.UMASS.EDU
- Iterated Query
 - * Contacted server replies with name of server to contact
 - * “I don’t know this name, but ask this server”

- DNS Name Resolution: Recursive Query

- Same situation as above
- Recursive Query
 - * Puts burden of name resolution on contacted name server
 - * Heavy load at upper levels of hierarchy?
- Caching, Updating DNS Records
 - Once (any) local DNS server learns mapping, it caches mapping
 - * Cache entries timeout (disappear) after some time (TTL)
 - * TLD servers typically cached in local name servers
 - Thus, root name servers not often visited
 - * Cached entries may be out-of-date (best effort name-to-address translation)
 - If name host changes IP address, may not be known Internet-wide until all TTLs expire
 - * Update/notify mechanisms proposed IETF standard
 - RFC 2136
- DNS Records
 - DNS: distributed database storing Resource Records (RR)

RR Format: (NAME, VALUE, TYPE, TTL)

 - TYPE=A
 - * NAME is hostname
 - * VALUE is IP address
 - TYPE=NS
 - * NAME is domain (e.g., foo.com)
 - * VALUE is hostname of authoritative name server for this domain
 - TYPE=CNAME
 - * NAME is alias name for some “canonical” (the real) name
 - WWW.IBM.COM is really SERVEREAST.BACKUP2.IBM.COM
 - * VALUE is canonical name
 - TYPE=MX
 - * NAME is a mail domain
 - * VALUE is hostname of mail server associated with NAME
- DNS Protocol Messages
 - DNS query and reply messages, both have same format:

- * Message header:
 - Identification — 16 bit number for query; reply to query uses same number
 - Flags — query or reply; recursion desired/available; if reply is authoritative
- Inserting Records into DNS
 - DNS registrar: commercial entity that verifies the uniqueness of a domain name
 - * Enters the domain name into the DNS database
 - * Collects a small fee from you for its services
 - There are many registrars that compete for customers
 - The ICAAN accredits the various registrars
 - Example: a new start-up “Network Utopia”
 - Register domain name NETWORKUTOPIA.COM at some DNS registrar (e.g., Network Solutions)
 - * Provide names, IP addresses of authoritative name server (primary and secondary)
 - * Registrar inserts NS and A type RRs into .com TLD server for each authoritative name server:

(NETWORKUTOPIA.COM, ADNS1.NETWORKUTOPIA.COM, NS)

(ADNS1.NETWORKUTOPIA.COM, 212.212.212.1, A)

- Create authoritative server locally with IP address 212.212.212.1 and enter:
 - * Type A record for your web server: (WWW.NETWORKUTOPIA.COM, “IP ADDRESS”, A)
 - * Type MX record for your mail server:

(NETWORKUTOPIA.COM, MAIL.NETWORKUTOPIA.COM, MX)

(MAIL.NETWORKUTOPIA.COM, ”IP ADDRESS”, A)

- Peer-To-Peer (P2P) Architecture
 - No always-on server
 - Arbitrary end systems directly communicate
 - Peers are intermittently connected and change IP addresses
 - Examples:
 - * File distribution (BitTorrent)

- * Streaming (KanKan)
 - * VoIP (Skype)
- P2P File Distribution: BitTorrent
 - File divided into 256kb chunks
 - Peers in torrent send/receive file chunks
 - Tracker: tracks peers participating in torrent
 - Torrent: group of peers exchanging chunks of a file
 - Peer joining torrent:
 - * Has no chunks, but will accumulate them over time from other peers
 - * Registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
 - While downloading, peer uploads chunks to other peers
 - Peer may change peers with whom it exchanges chunks
 - Churn: peers may come and go
 - Requesting Chunks:
 - * At any given time, different peers have different subsets of file chunks
 - * Periodically, a certain peer asks each peer for list of chunks they have
 - * The peer then requests missing chunks from peers, rarest first
 - Sending Chunks: tit-for-tat
 - * A peer sends chunks to those four peers currently sending them chunks at highest rate
 - Other peers are choked by the peer (do not receive chunks)
 - Re-evaluate top 4 every 10 seconds
 - * Every 30 seconds randomly “unchoke” another peer
 - * Newly chosen peer may join top 4
- Video Streaming and CDNs: Context
 - Stream video traffic: major consumer of internet bandwidth
 - * Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic
 - Challenge: scale — how to reach about 1B users?
 - * Single mega-video server won’t work
 - Challenge heterogeneity
 - * Different users have different capabilities
 - Solution: distributed, application-level infrastructure

- Multimedia: Video
 - Video: sequence of images displayed at constant rate
 - * Like 24 frames per second
 - Digital image: array of pixels
 - * Each pixel represented by bits
 - Coding: use redundancy within and between images to decrease # used to encode image
 - * Spatial (within image) — Instead of sending N values of same color (all purple) send only two values: color value (purple) and number of repeated values (N)
 - * Temporal (from one image to the next) — Instead of sending complete frame at $i + 1$, send only differences from frame i
 - CBR (Constant bit rate): video encoding rate fixed
 - VBR (Variable bit rate): video encoding rate changes as amount of spatial, temporal coding changes
 - Examples:
 - * MPEG1 (CD-ROM) 1.5 Mbps
 - * MPEG2 (DVD) 3-6 Mbps
 - * MPEG4 (Used in Internet) 64kbps-12Mbps
 - Main challenges:
 - * Server-to-client rate will vary over time, with changing network congestion levels (in house, in access network, in network core, at video server)
 - * Packet loss and delay due to congestion will delay playout, or result in poor video quality
- Streaming stored video
 - Continuous playout constraint: once client playout begins, playback must match original timing
 - * But network delays are variable (jitter), so we will need client-side buffer to match playout requirements
 - Client-side buffering and playout delay: compensate for network-added delay and jitter
- Streaming Multimedia: DASH
 - DASH: Dynamic, Adaptive Streaming over HTTP
 - Server:

- * Divides video file into multiple chunks
- * Each chunk stored, encoded at different rates
- * Manifest file: provides URLs for different chunks
- Client:
 - * Periodically measures server-to-client transmission rate
 - * Consulting manifest, requests one chunk at a time
 - Chooses maximum coding rate sustainable given current transmission rate
 - Can choose different coding rates at different points in time (depending on available transmission rate at time)
 - * “Intelligence” at client: client determines:
 - When to request chunk (so that buffer starvation, or overflow, does not occur)
 - What encoding rate to request (higher quality when more transmission rate available)
 - Where to request chunk (can request from URL server that is “close” to client or has high available transmission rate)
- Content Distribution Networks (CDNs)
 - Challenge: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
 - Option 1: single, large “mega-server”
 - * Single point of failure
 - * Point of network congestion
 - * Long path to distant clients
 - * Multiple copies of video sent over outgoing link
 - * Doesn’t scale!
 - Option 2: store/server multiple copies of videos at multiple geographically distributed sites (CDN)
 - * Enter deep: push CDN servers deep into many networks of access ISPs
 - Close to users
 - Akamai, 240000 servers deployed in more than 120 countries
 - * Bring home: smaller number (10’s) of larger clusters in POPs (Points of Presence) near (but not within) access ISPs
 - Typically, POPs are Internet Exchange Points (IXPs)
 - Used by Limelight
 - CDN: stores copies of content at CDN nodes
 - Subscriber requests content from CDN

- * Directed to nearby copy, retrieves content
- * May choose different copy if network path congested
- Over-the-top (OTT)
 - * OTT challenges: coping with a congested Internet
 - From which CDN node to retrieve content?
 - Viewer behavior in presence of congestion?
 - What content to place in which CDN node?