

The Transport Layer

Michael Brodskiy

Professor: E. Bernal Mor

October 11, 2023

- Transport Services and Protocols
 - Provide logical communication between application processes running on different hosts
 - Transport protocols actions in end systems:
 - * Sender: breaks application messages into segments, passes to Network layer
 - * Receiver: reassembles segments into messages, passes to Application layer
 - Two transport protocols available to internet applications
 1. TCP
 2. UDP
- Transport vs. Network Layer
 - Network layer: logical communication between two hosts
 - Transport layer: logical communication between processes
 - * Relies on, enhances, network layer services
- Two Internet Transport Protocols
 - TCP: Transmission Control Protocol
 - * Reliable, in-order delivery
 - * Congestion Control
 - * Flow Control
 - * Connection set-up
 - UDP: User Datagram Protocol
 - * Unreliable, unordered delivery
 - * No-frills extension of “best-effort” IP

- Services not available:
 - * Delay guarantees
 - * Throughput guarantees
- Multiplexing/Demultiplexing
 - Multiplexing at sender: Handle data from multiple sockets, add transport header (later used for demultiplexing)
 - How demultiplexing works
 - * Host receives IP packets
 - Each packet has source IP address, destination IP address
 - Each packet carries one transport-layer segment
 - Each segment has source, destination port number
 - * Host uses IP address and port numbers to direct segment to appropriate socket
 - Connectionless Demultiplexing
 - * Create a socket in the client, the Transport layer automatically assigns a host-local port number to the socket
 - * When data is sent into UDP socket, must specify
 - Destination IP address
 - Destination port number
 - * When a host receives UDP segment, the Transport layer:
 - Checks destination port number in segment
 - Directs UDP segment to socket with that port number
 - * IP datagrams with same destination port number but different source IP addresses and/or source port numbers will be directed to same socket at destination
 - Connection-Oriented Demultiplexing
 - * TCP socket identified by 4-tuple:
 - Source IP address
 - Source port number
 - Destination IP address
 - Destination port number
 - * Demultiplexing receiver uses all four values to direct segment to appropriate socket
 - * A server may support simultaneous TCP sockets:
 - Each socket identified by its own 4-tuple
 - Each socket associated with a different connecting client

- * Note: the TCP server has a welcoming socket
 - Each time a client initiates a TCP connection to the server, a new socket is created for this connection
 - To support n simultaneous connections, the server would need $n + 1$ sockets
- Connectionless Transport: UDP
 - “No frills”, “bare bones” Internet transport protocol
 - “Best effort” service, UDP segments may be:
 - * Lost
 - * Delivered out-of-order to application
 - Connectionless:
 - * No handshaking between UDP sender, receiver
 - * Each UDP segment handled independently of others
 - Why is there a UDP?
 - * No connection establishment (which can add RTT delay)
 - * Simple: no connection state at sender, receiver
 - * Small header size
 - * No congestion control
 - UDP can blast away as fast as desired
 - It can function in the face of congestion
 - UDP used in:
 - * Streaming multimedia apps (loss tolerant, rate sensitive)
 - * DNS
 - * HTTP/3
 - If reliable transfer or other services needed over UDP (like in HTTP/3)
 - * Add needed reliability at application layer
 - * Add congestion control at application layer
- Internet Checksum
 - Goal: detect “errors” (*e.g.* flipped bits) in received segment — optional if UDP segment is encapsulated in IPv4 packet: carries all-zeros if unused
 - Sender:
 - * Treat segment contents, including UDP header fields and IP addresses, as sequence of 16-bit words
 - * Checksum: ones complement of the ones complement sum of all 16-bit words

- * Sender puts checksum value into UDP checksum field
- Receiver — error detected?
 - * All 16-bit words are added, including checksum
 - * Result = 1111111111111111 → no error
 - But maybe errors, nonetheless?
 - * Result \neq 1111111111111111 → Error
 - Do not recover from the error: discard segment or pass damaged data with a warning
- Principles of Reliable Data Transfer
 - Consider only unidirectional data transfer
 - But control information will flow in both directions
 - Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- Components of Error Recovery
 - Checksum to detect bit errors
 - Acknowledgments (ACKs): receiver explicitly tells sender the packet received is OK
 - Negative Acknowledgments (NAKs): receiver explicitly tells sender that pack had errors
 - Sender retransmits packet on receipt of NAK
 - What happens if ACK/NAK is corrupted?
 - * Sender doesn't know what happened at receiver
 - * Can't just retransmit, possible duplicate
 - Sender:
 - * Sequence number added to packet
 - * Two sequence numbers will suffice
 - * Must check if received ACK/NAK corrupted
 - * Sender must “remember” whether last sent packet had sequence number of 0 or 1
 - Receiver:
 - * Must check if received packet is duplicate
 - Receiver must remember whether 0 or 1 is expected packet sequence number
 - * If duplicate, send ACK to force sender to move on
 - * Note: receiver can not know if its last NAK/ACK received ok at sender

- A similar protocol without NAKs may be developed:
 - * Instead of NAK, receiver sends ACK for last packet received OK
 - Receiver must explicitly include in the ACK the sequence number of packet being ACKed
 - * Duplicate ACK at sender results in same action as NAK: retransmit current packet
 - * TCP uses this approach to be NAK-free
- This protocol may be further developed:
 - * Approach: sender waits “reasonable” amount of time for ACK
 - * Only retransmit if no expected ACK received in this time
 - * If packet (or ACK) just delayed (not lost):
 - Retransmission will be duplicate, but sequence number already handles this
 - Receiver must specify sequence number of packet being ACKed
 - * Requires countdown timer to interrupt after “reasonable” amount of time (timeout)

- Utility

- For a “Stop and Wait” approach, the utility can be defined as:

$$U_{sender} = \frac{L/R}{RTT + (L/R)}$$

- Thus, the performance of stop and wait would not be good if the RTT is high
- Stop and wait limits performance of underlying infrastructure

- Pipelined Protocols

- Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged packets
 - * Range of sequence numbers must be increased k -bit sequence number in packet header
 - * Buffering at sender and/or receiver
- Two generic forms of pipelined protocols: Go-Back-N (GBN) and selective repeat