

# The Transport Layer

Michael Brodskiy

Professor: E. Bernal Mor

October 12, 2023

- Transport Services and Protocols
  - Provide logical communication between application processes running on different hosts
  - Transport protocols actions in end systems:
    - \* Sender: breaks application messages into segments, passes to Network layer
    - \* Receiver: reassembles segments into messages, passes to Application layer
  - Two transport protocols available to internet applications
    1. TCP
    2. UDP
- Transport vs. Network Layer
  - Network layer: logical communication between two hosts
  - Transport layer: logical communication between processes
    - \* Relies on, enhances, network layer services
- Two Internet Transport Protocols
  - TCP: Transmission Control Protocol
    - \* Reliable, in-order delivery
    - \* Congestion Control
    - \* Flow Control
    - \* Connection set-up
  - UDP: User Datagram Protocol
    - \* Unreliable, unordered delivery
    - \* No-frills extension of “best-effort” IP

- Services not available:
  - \* Delay guarantees
  - \* Throughput guarantees
- Multiplexing/Demultiplexing
  - Multiplexing at sender: Handle data from multiple sockets, add transport header (later used for demultiplexing)
  - How demultiplexing works
    - \* Host receives IP packets
      - Each packet has source IP address, destination IP address
      - Each packet carries one transport-layer segment
      - Each segment has source, destination port number
    - \* Host uses IP address and port numbers to direct segment to appropriate socket
  - Connectionless Demultiplexing
    - \* Create a socket in the client, the Transport layer automatically assigns a host-local port number to the socket
    - \* When data is sent into UDP socket, must specify
      - Destination IP address
      - Destination port number
    - \* When a host receives UDP segment, the Transport layer:
      - Checks destination port number in segment
      - Directs UDP segment to socket with that port number
    - \* IP datagrams with same destination port number but different source IP addresses and/or source port numbers will be directed to same socket at destination
  - Connection-Oriented Demultiplexing
    - \* TCP socket identified by 4-tuple:
      - Source IP address
      - Source port number
      - Destination IP address
      - Destination port number
    - \* Demultiplexing receiver uses all four values to direct segment to appropriate socket
    - \* A server may support simultaneous TCP sockets:
      - Each socket identified by its own 4-tuple
      - Each socket associated with a different connecting client

- \* Note: the TCP server has a welcoming socket
  - Each time a client initiates a TCP connection to the server, a new socket is created for this connection
  - To support  $n$  simultaneous connections, the server would need  $n + 1$  sockets
- Connectionless Transport: UDP
  - “No frills”, “bare bones” Internet transport protocol
  - “Best effort” service, UDP segments may be:
    - \* Lost
    - \* Delivered out-of-order to application
  - Connectionless:
    - \* No handshaking between UDP sender, receiver
    - \* Each UDP segment handled independently of others
  - Why is there a UDP?
    - \* No connection establishment (which can add RTT delay)
    - \* Simple: no connection state at sender, receiver
    - \* Small header size
    - \* No congestion control
      - UDP can blast away as fast as desired
      - It can function in the face of congestion
  - UDP used in:
    - \* Streaming multimedia apps (loss tolerant, rate sensitive)
    - \* DNS
    - \* HTTP/3
  - If reliable transfer or other services needed over UDP (like in HTTP/3)
    - \* Add needed reliability at application layer
    - \* Add congestion control at application layer
- Internet Checksum
  - Goal: detect “errors” (*e.g.* flipped bits) in received segment — optional if UDP segment is encapsulated in IPv4 packet: carries all-zeros if unused
  - Sender:
    - \* Treat segment contents, including UDP header fields and IP addresses, as sequence of 16-bit words
    - \* Checksum: ones complement of the ones complement sum of all 16-bit words

- \* Sender puts checksum value into UDP checksum field
- Receiver — error detected?
  - \* All 16-bit words are added, including checksum
  - \* Result = 1111111111111111 → no error
    - But maybe errors, nonetheless?
  - \* Result  $\neq$  1111111111111111 → Error
    - Do not recover from the error: discard segment or pass damaged data with a warning
- Principles of Reliable Data Transfer
  - Consider only unidirectional data transfer
  - But control information will flow in both directions
  - Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- Components of Error Recovery
  - Checksum to detect bit errors
  - Acknowledgments (ACKs): receiver explicitly tells sender the packet received is OK
  - Negative Acknowledgments (NAKs): receiver explicitly tells sender that pack had errors
  - Sender retransmits packet on receipt of NAK
  - What happens if ACK/NAK is corrupted?
    - \* Sender doesn't know what happened at receiver
    - \* Can't just retransmit, possible duplicate
  - Sender:
    - \* Sequence number added to packet
    - \* Two sequence numbers will suffice
    - \* Must check if received ACK/NAK corrupted
    - \* Sender must “remember” whether last sent packet had sequence number of 0 or 1
  - Receiver:
    - \* Must check if received packet is duplicate
      - Receiver must remember whether 0 or 1 is expected packet sequence number
    - \* If duplicate, send ACK to force sender to move on
    - \* Note: receiver can not know if its last NAK/ACK received ok at sender

- A similar protocol without NAKs may be developed:
  - \* Instead of NAK, receiver sends ACK for last packet received OK
    - Receiver must explicitly include in the ACK the sequence number of packet being ACKed
  - \* Duplicate ACK at sender results in same action as NAK: retransmit current packet
  - \* TCP uses this approach to be NAK-free
- This protocol may be further developed:
  - \* Approach: sender waits “reasonable” amount of time for ACK
  - \* Only retransmit if no expected ACK received in this time
  - \* If packet (or ACK) just delayed (not lost):
    - Retransmission will be duplicate, but sequence number already handles this
    - Receiver must specify sequence number of packet being ACKed
  - \* Requires countdown timer to interrupt after “reasonable” amount of time (timeout)

- Utility

- For a “Stop and Wait” approach, the utility can be defined as:

$$U_{sender} = \frac{L/R}{RTT + (L/R)}$$

- Thus, the performance of stop and wait would not be good if the RTT is high
- Stop and wait limits performance of underlying infrastructure

- Pipelined Protocols

- Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged packets
  - \* Range of sequence numbers must be increased  $k$ -bit sequence number in packet header
  - \* Buffering at sender and/or receiver
- Two generic forms of pipelined protocols: Go-Back-N (GBN) and selective repeat

- Go-Back-N

- A “sliding window” protocol
- Sender window:  $N$  consecutive sequence numbers allowed for sent, unacknowledged packets
  - \* Sender window size is  $N$

- Cumulative ACK:  $\text{ACK}(n) \rightarrow$  Acknowledges all packets up to  $N$ , including sequence number  $n$ 
  - \* On receiving  $\text{ACK}(n) \rightarrow$  move window forward to begin at  $n + 1$
- Selective Repeat
  - Sender Window
    - \*  $N$  consecutive sequence numbers
    - \* Limits sequence numbers of sent, unacknowledged packets
    - \* Sender window size is  $N$
  - Sender maintains a timer for each unACKed packet
  - Sender times-out and retransmits individually unACKed packets
  - Receiver Window
    - \*  $N$  consecutive sequence number
    - \* Limits sequence numbers of received packets that are accepted
    - \* Receiver window size is  $N$
  - Buffers packets, as needed for eventual in-order delivery to upper layer
  - Receiver individually acknowledges all correctly received packets
- Connection-Oriented Transport: TCP
  - Point-to-point: one sender, one receiver
  - Reliable, in-order byte stream:
    - \* No “message boundaries”
  - Full duplex data:
    - \* Bi-directional data flow in same connection
    - \* MSS: maximum segment size
  - Cumulative ACKs
  - Pipelining:
    - \* TCP congestion and flow control set the window size
  - Connection-oriented
    - \* Handshaking (exchange of control messages) initializes sender and receiver state before data exchange
  - Flow control
    - \* Sender will not overwhelm receiver